KHUSHI AGRAWAL

(623)212-5340 • khushiagra011@gmail.com • linkedin.com/in/khushi-agrawal-83aa0a202/
Graduation: May 2026 | Available for full-time roles starting Summer/Fall 2026

EDUCATION

Master of Science, Computer Science **Aug 2024 - May 2026**
Arizona State University, Tempe, AZ (**CGPA - 4.0/4.0)**
Relevant Courses: Data Processing at Scale, Data Mining, Semantic Web Mining, Statistical Machine Learning

Bachelor of Technology, Computer Science and Engineering **Jul 2019 – Jul 2023**
Vellore Institute of Technology, Vellore, India (**CGPA – 3.56/4.0)**
Relevant Courses: Machine Learning, Data Warehousing, Image Processing, Artificial Intelligence, Cybersecurity

TECHNICAL SKILLS

Programming Languages: Python, JavaScript, Java, SQL, Apex (basic), C/C++
AI/LLMs: Prompt Engineering, LLM Integration (OpenAI/Anthropic APIs), Finetuning, Agent Orchestration, TensorFlow, PyTorch
Data & Platforms: Salesforce (Admin + Platform basics), Snowflake, Databricks, Kafka, PySpark, Hadoop, GCP, AWS, Docker
Data Engineering: ETL/ELT Pipelines, Data Modeling, Integration, Workflow Automation
Tools: Git, Jira, REST APIs, MLflow, Postman, Linux, Docker Compose

PROFESSIONAL EXPERIENCE

DruvStar, Tempe, USA: Data Analyst Engineer Intern **May 2025 – Present**
- Delivered end-to-end AI/LLM-powered pipelines and integrated them using Salesforce-style data modeling principles, orchestration patterns, and API-driven workflows.
- Built containerized microservices using Docker, Kafka, and Python, enabling scalable offline inference with 3× throughput and seamless integration into production systems.
- Implemented hands-on debugging, telemetry instrumentation, and pipeline optimization to remove technical blockers across ML inference, data flow, and orchestration layers.
- Partnered directly with security and DevOps teams to integrate robust data models, enforce privacy constraints, and maintain high system reliability.

Accenture, Hyderabad, India: Advanced App Engineering Analyst **Oct 2023 – Aug 2024**
- Built large-scale PySpark pipelines processing 5TB/day, enabling real-time analytics for operational and ML-driven use cases.
- Designed reusable data validation modules, reducing data inconsistencies by 70% and improving downstream model performance.
- Co-led experimentation frameworks supporting A/B testing and rapid iteration cycles, resulting in 12% lift in model ROI.
- Resolved complex data integration issues across distributed systems through root-cause debugging and schema optimization.

Ataloud Technologies, Mumbai, India: Full Stack Developer Intern **Sep 2022 – Mar 2023**
- Integrated 40+ REST APIs, ensuring smooth transfer of structured and unstructured data across services and improving reliability of business workflows.
- Enhanced frontend data rendering in React/Next.js for 10,000+ users, improving accuracy and performance of analytics interfaces.
- Diagnosed and resolved 77% of high-priority API data issues through hands-on debugging and close collaboration with backend teams.

ACADEMIC PROJECTS

LLMind, Tempe, AZ **Spring 2025**
- Designed an LLM-powered analytical agent for multimodal sentiment forecasting across 1M+ financial tweets, integrating prompt engineering, API-based LLM reasoning, and ML pipelines.
- Built data transformation workflows and MLflow tracking to accelerate experimentation by 50%.
- Conducted deep error analysis and iterative refinement to optimize model reasoning and contextual understanding.

AI-Driven Forecasting of Dengue Outbreaks, Tempe, AZ **Fall 2024**
- Built ARIMA–LSTM hybrid forecasting models and automated data preparation workflows for geospatial public-health datasets.
- Designed reproducible pipelines for model tuning, validation, and versioning to support rapid experimentation.